

+

# Machine Learning and Data Mining

## Multi-layer Perceptrons & Neural Networks: Basics

Prof. Alexander Ihler



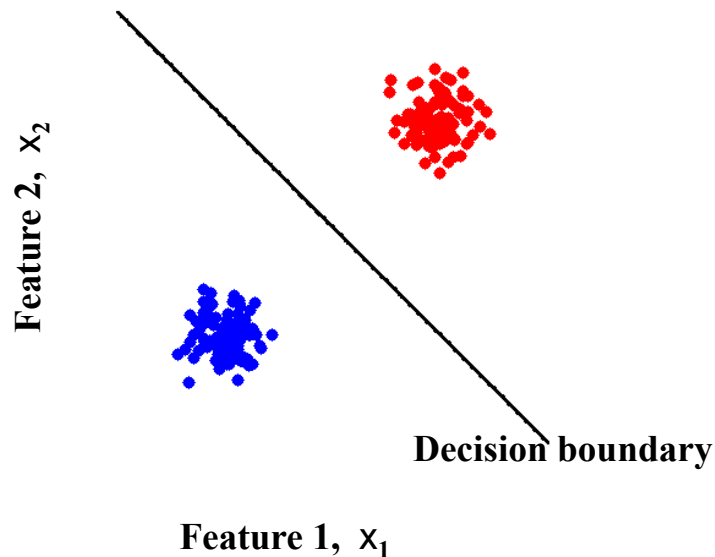
**BREN:ICS**  
INFORMATION AND COMPUTER SCIENCES

UNIVERSITY of CALIFORNIA  IRVINE

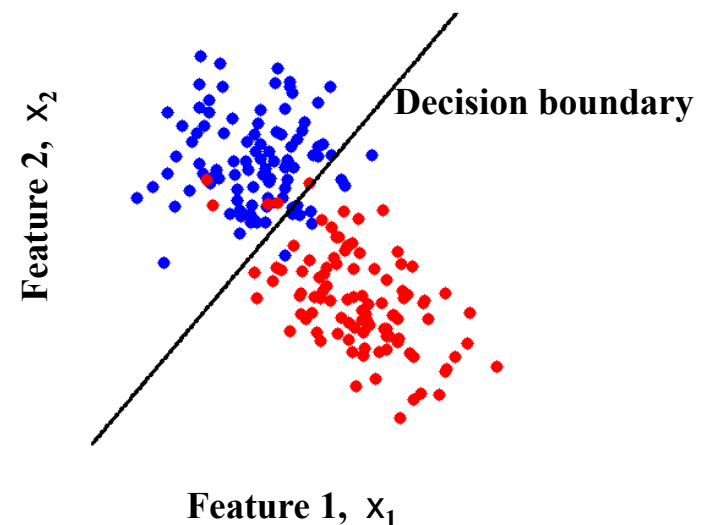
# Linear classifiers (perceptrons)

- Linear Classifiers
  - a linear classifier is a mapping which partitions feature space using a linear function (a straight line, or a hyperplane)
  - separates the two classes using a straight line in feature space
  - in 2 dimensions the decision boundary is a straight line

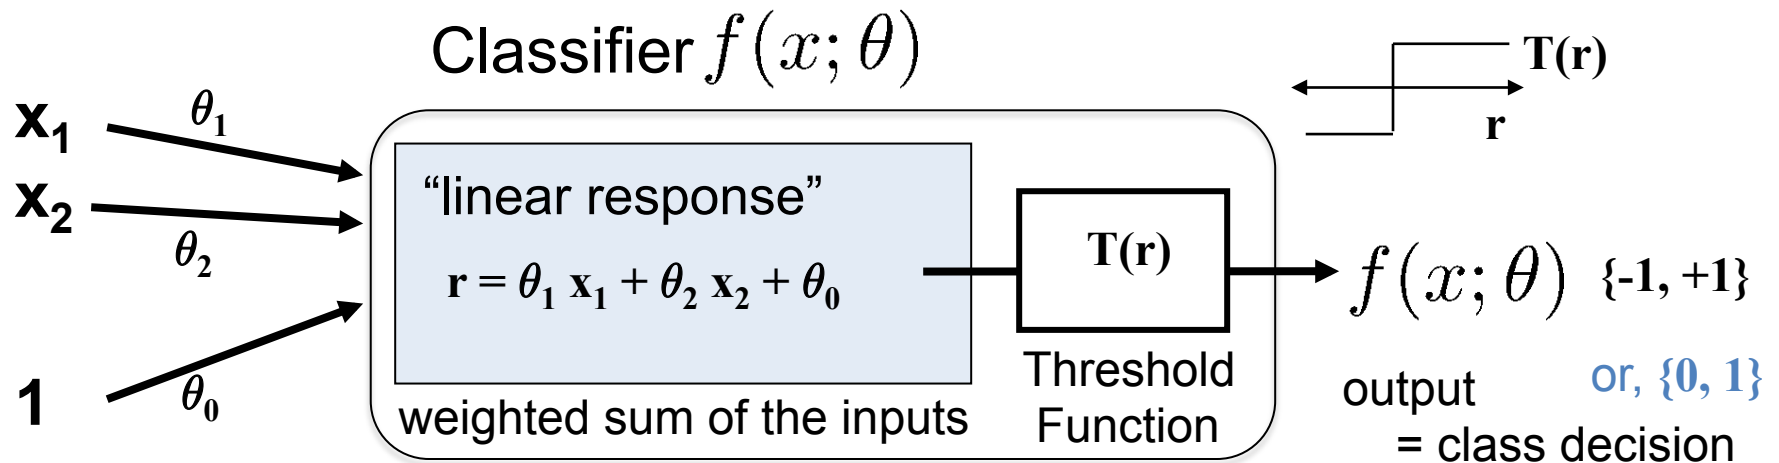
Linearly separable data



Linearly non-separable data



# Perceptron Classifier (2 features)

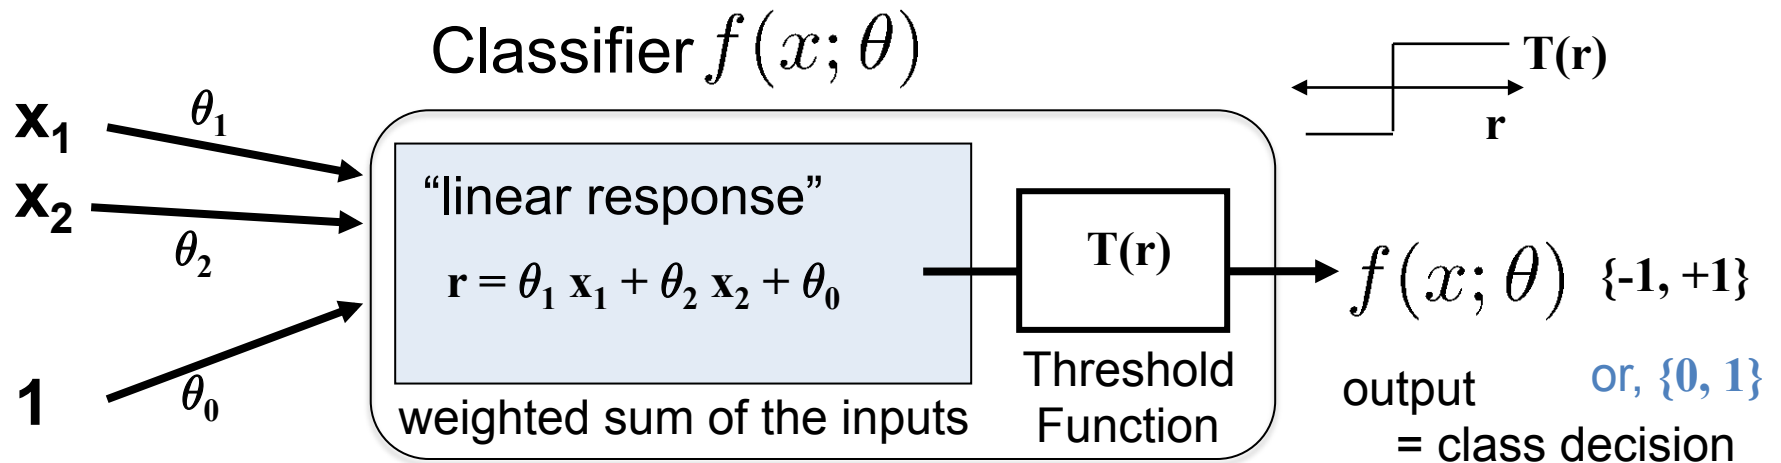


```
r = X.dot( theta.T );           # compute linear response
Yhat = 2*(r > 0)-1             # "sign": predict +1 / -1
```

Decision Boundary at  $r(x) = 0$

Solve:  $X_2 = -w_1/w_2 X_1 - w_0/w_2$  (Line)

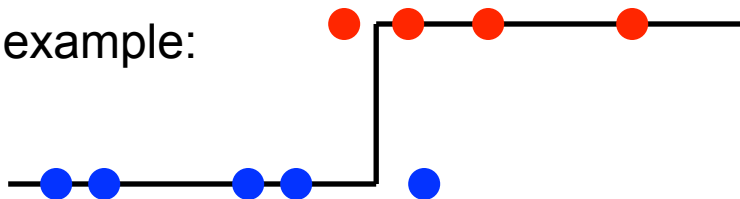
# Perceptron Classifier (2 features)



```

r = X.dot( theta.T );           # compute linear response
Yhat = 2*(r > 0)-1             # "sign": predict +1 / -1
    
```

1D example:



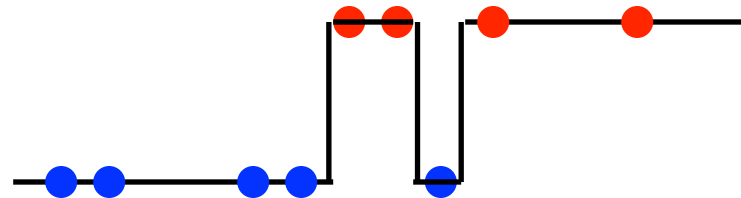
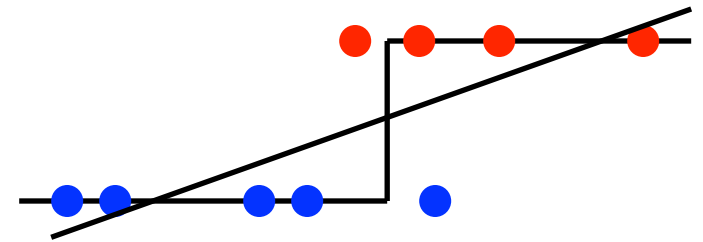
$$T(r) = -1 \text{ if } r < 0$$

$$T(r) = +1 \text{ if } r > 0$$

Decision boundary = “x such that  $T( w_1 x + w_0 )$  transitions”

# Features and perceptrons

- Recall the role of features
  - We can create extra features that allow more complex decision boundaries
  - Linear classifiers
  - Features  $[1, x]$ 
    - Decision rule:  $T(ax+b) = ax + b >/< 0$
    - Boundary  $ax+b=0 \Rightarrow$  point
  - Features  $[1, x, x^2]$ 
    - Decision rule  $T(ax^2+bx+c)$
    - Boundary  $ax^2+bx+c = 0 = ?$
  - What features can produce this decision rule?

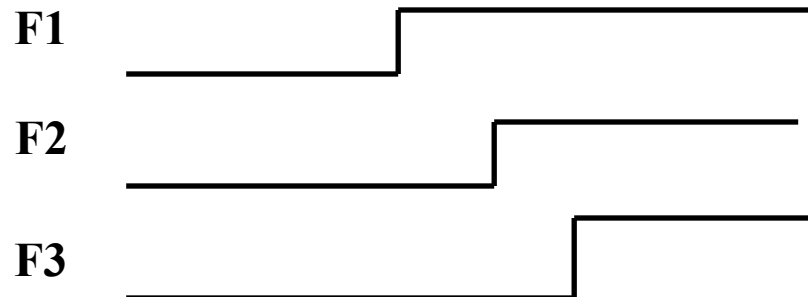


# Features and perceptrons

- Recall the role of features
  - We can create extra features that allow more complex decision boundaries
  - For example, polynomial features

$$\Phi(x) = [1 \ x \ x^2 \ x^3 \dots]$$

- What other kinds of features could we choose?
  - Step functions?



**Linear function of features**

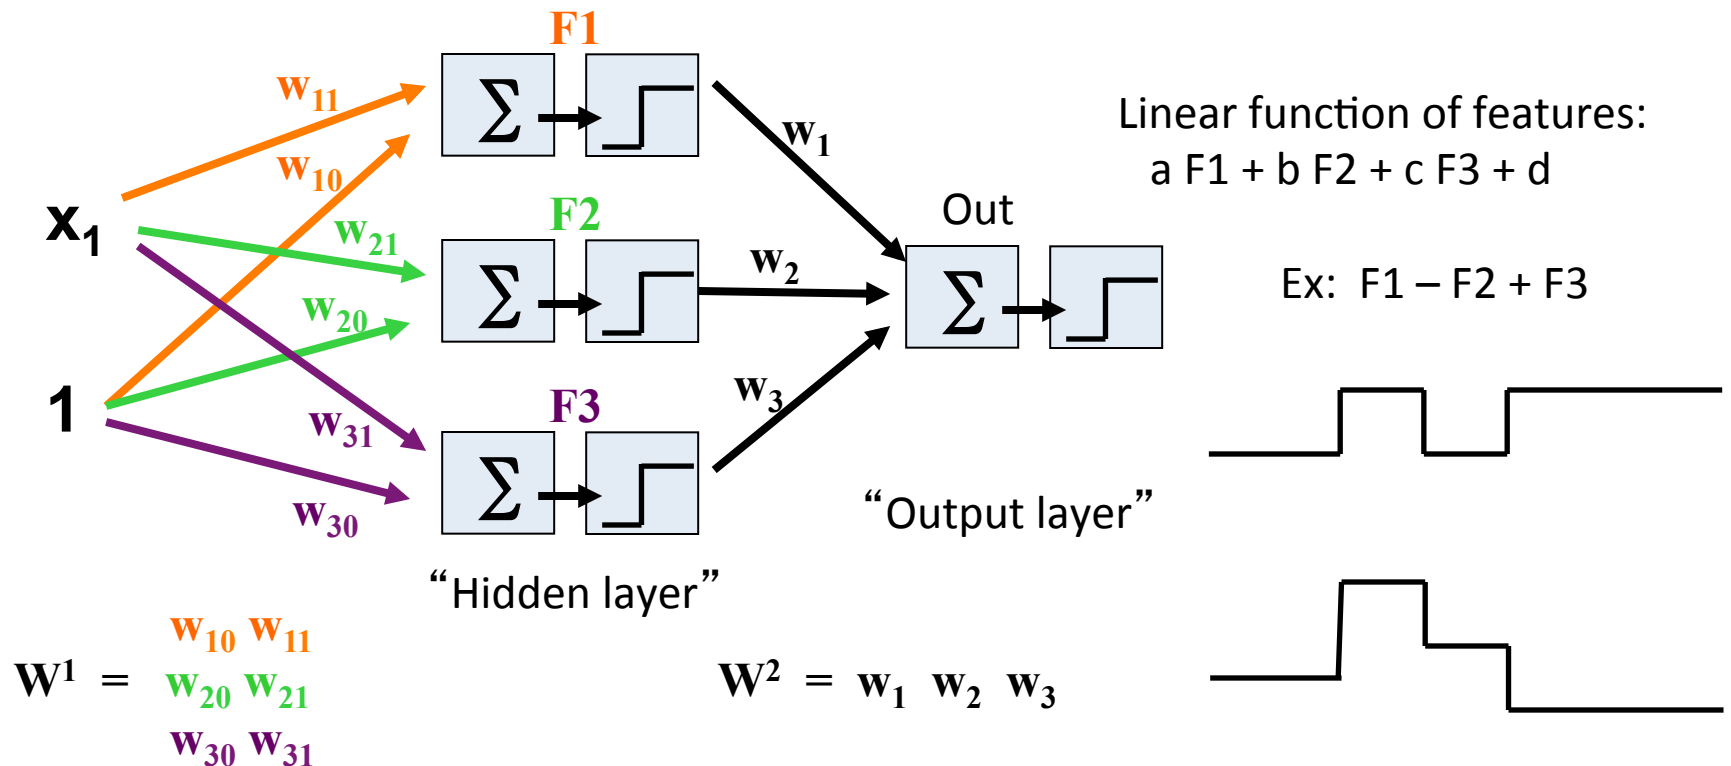
$$a F1 + b F2 + c F3 + d$$

**Ex:  $F1 - F2 + F3$**



# Multi-layer perceptron model

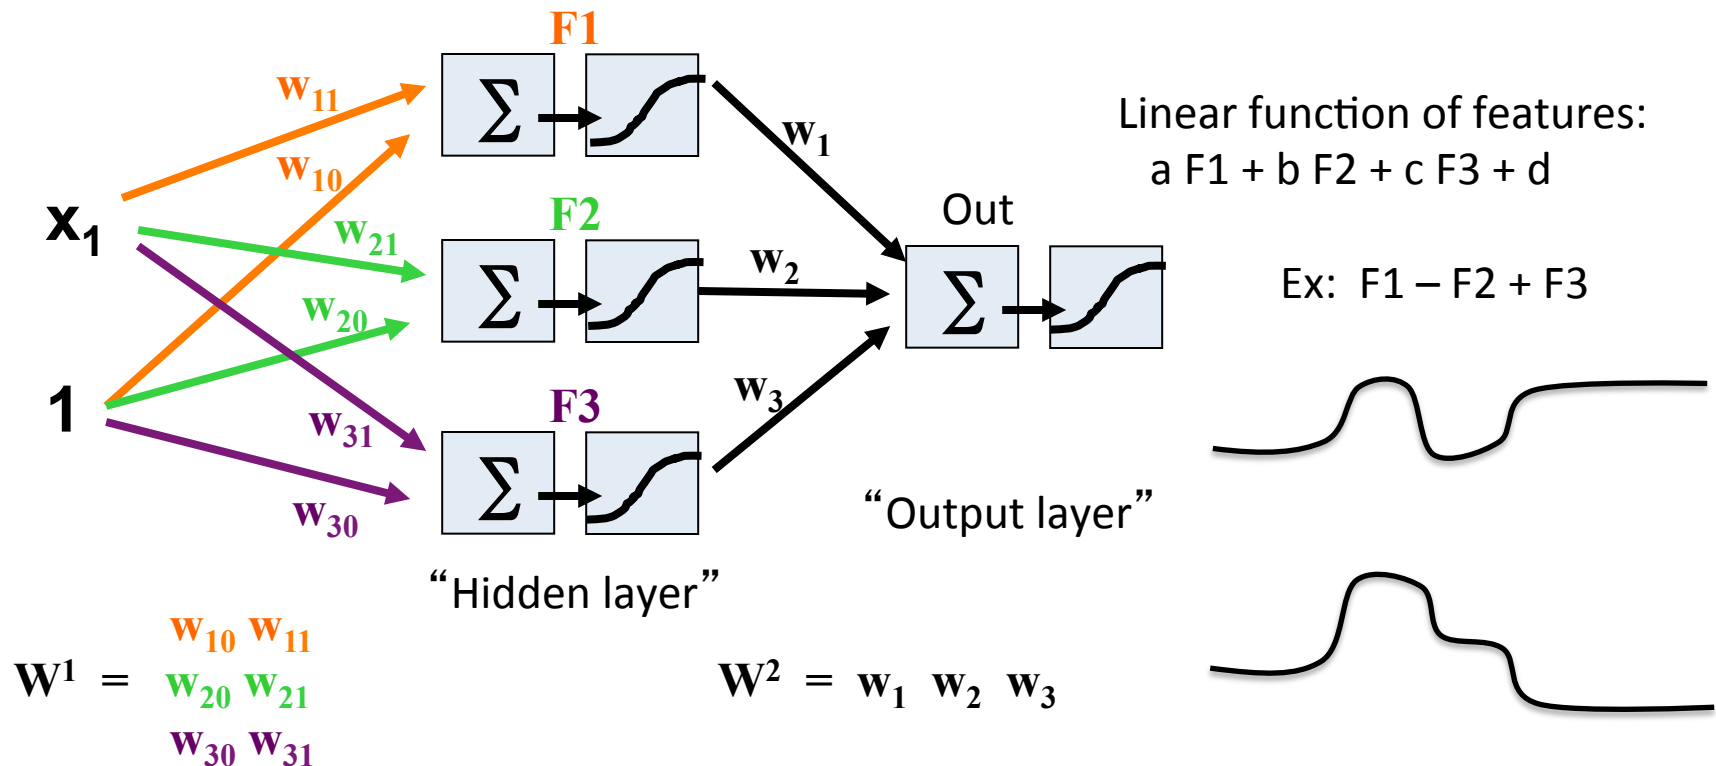
- Step functions are just perceptrons!
  - “Features” are outputs of a perceptron
  - Combination of features output of another



# Multi-layer perceptron model

- Step functions are just perceptrons!
  - “Features” are outputs of a perceptron
  - Combination of features output of another

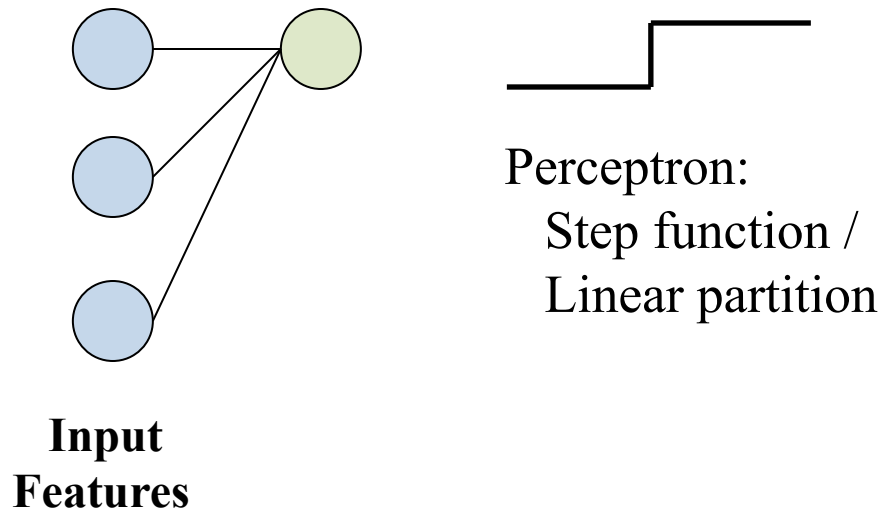
Regression version:  
Remove activation  
function from output





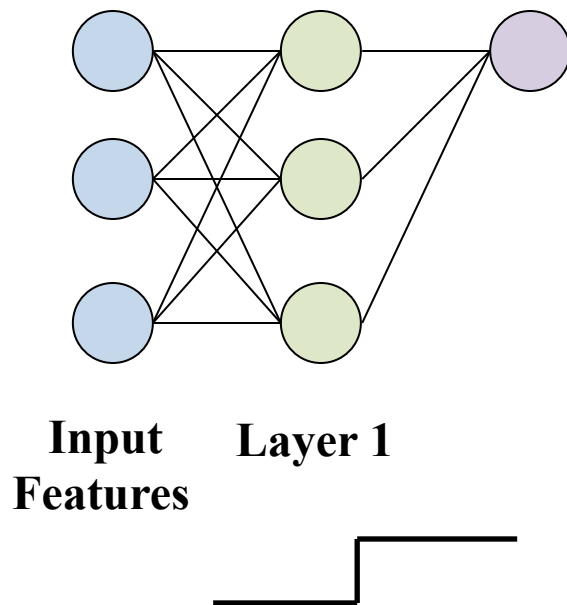
# Features of MLPs

- Simple building blocks
  - Each element is just a perceptron  $f'$  n
- Can build upwards



# Features of MLPs

- Simple building blocks
  - Each element is just a perceptron  $f'$  n
- Can build upwards



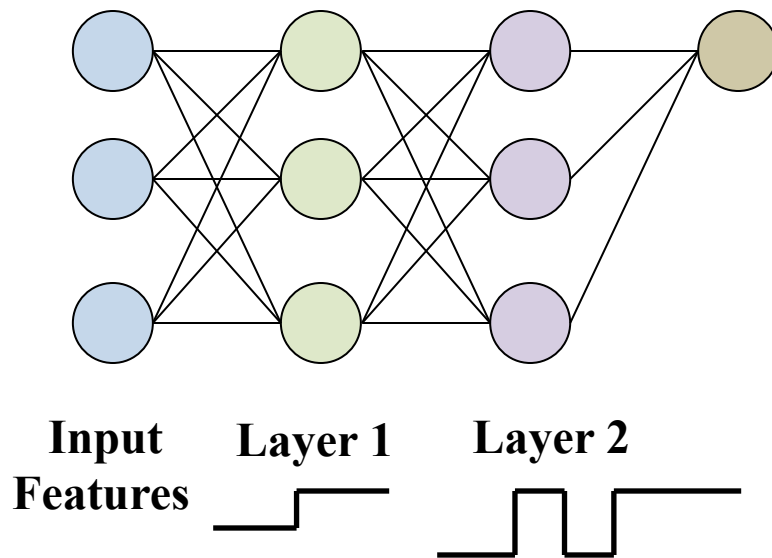
2-layer:

“Features” are now partitions

All linear combinations of those partitions

# Features of MLPs

- Simple building blocks
  - Each element is just a perceptron  $f'$  n
- Can build upwards

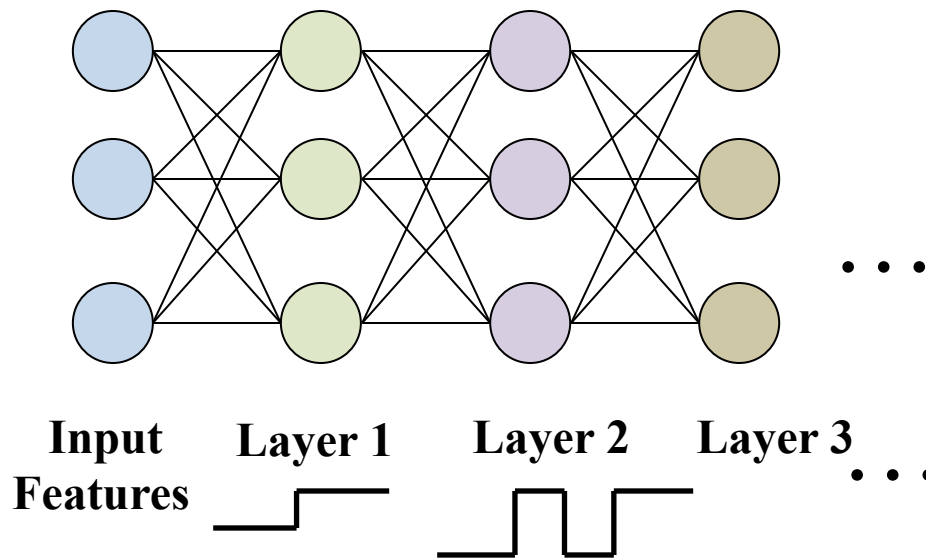


3-layer:

“Features” are now complex functions  
Output any linear combination of those

# Features of MLPs

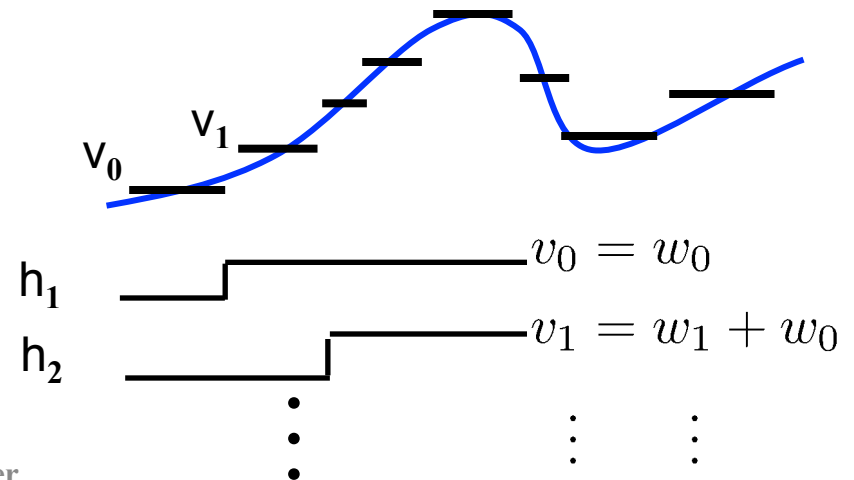
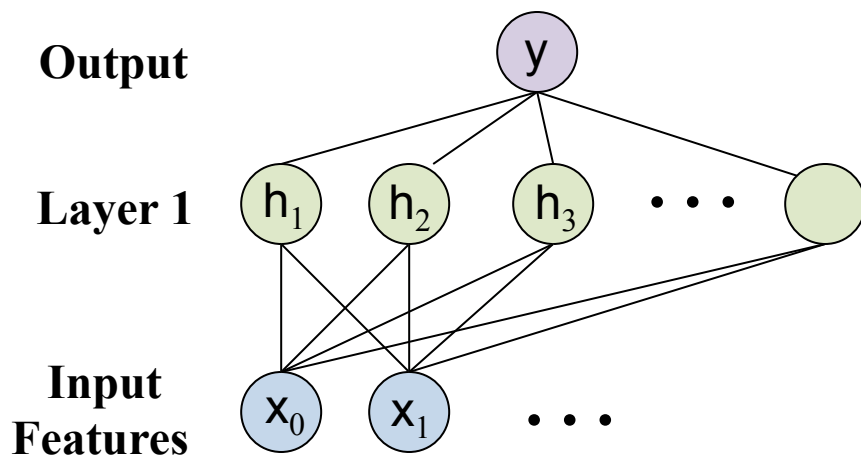
- Simple building blocks
  - Each element is just a perceptron  $f'$  n
- Can build upwards



Current research:  
“Deep” architectures  
(many layers)

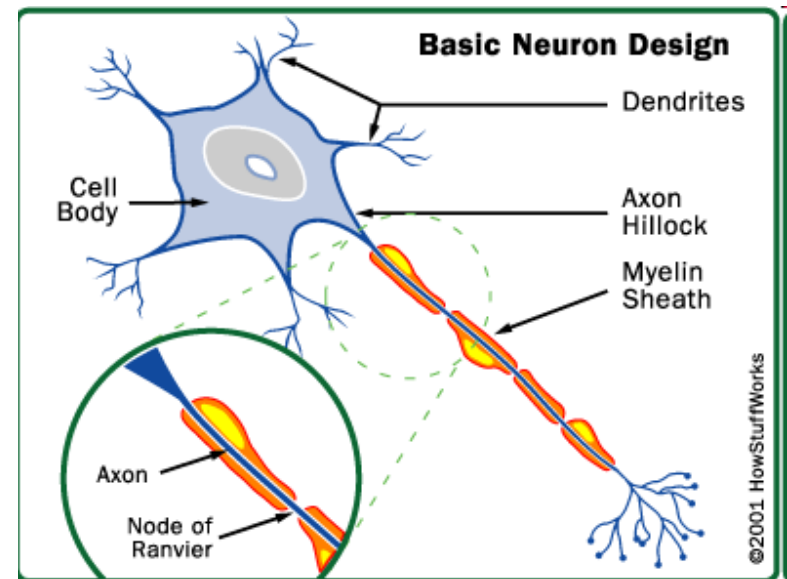
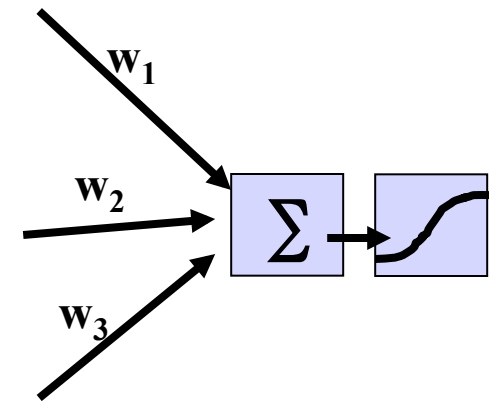
# Features of MLPs

- Simple building blocks
  - Each element is just a perceptron function
- Can build upwards
- Flexible function approximation
  - Approximate arbitrary functions with enough hidden nodes



# Neural networks

- Another term for MLPs
- Biological motivation
- Neurons
  - “Simple” cells
  - Dendrites sense charge
  - Cell weighs inputs
  - “Fires” axon



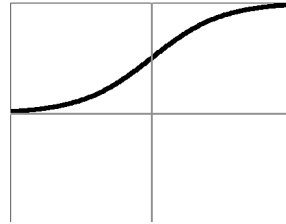
(c) Alexander Ihler

**“How stuff works: the brain”**

# Activation functions

Logistic

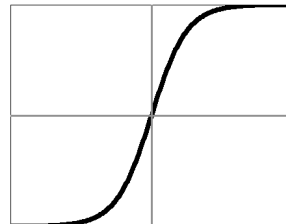
$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$



$$\frac{\partial \sigma}{\partial z}(z) = \sigma(z)(1 - \sigma(z))$$

Hyperbolic  
Tangent

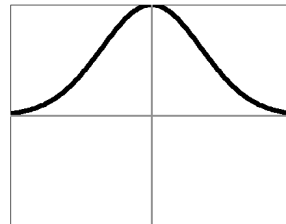
$$\sigma(z) = \frac{1 - \exp(-2z)}{1 + \exp(-2z)}$$



$$\frac{\partial \sigma}{\partial z}(z) = 1 - (\sigma(z))^2$$

Gaussian

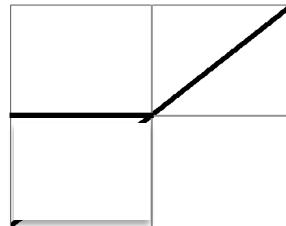
$$\sigma(z) = \exp(-z^2/2)$$



$$\frac{\partial \sigma}{\partial z}(z) = -z\sigma(z)$$

ReLU  
(rectified linear)

$$\sigma(z) = \max(0, z)$$



$$\frac{\partial \sigma}{\partial z}(z) = \mathbb{1}[z > 0]$$

Linear

$$\sigma(z) = z$$

and many others...

# Feed-forward networks

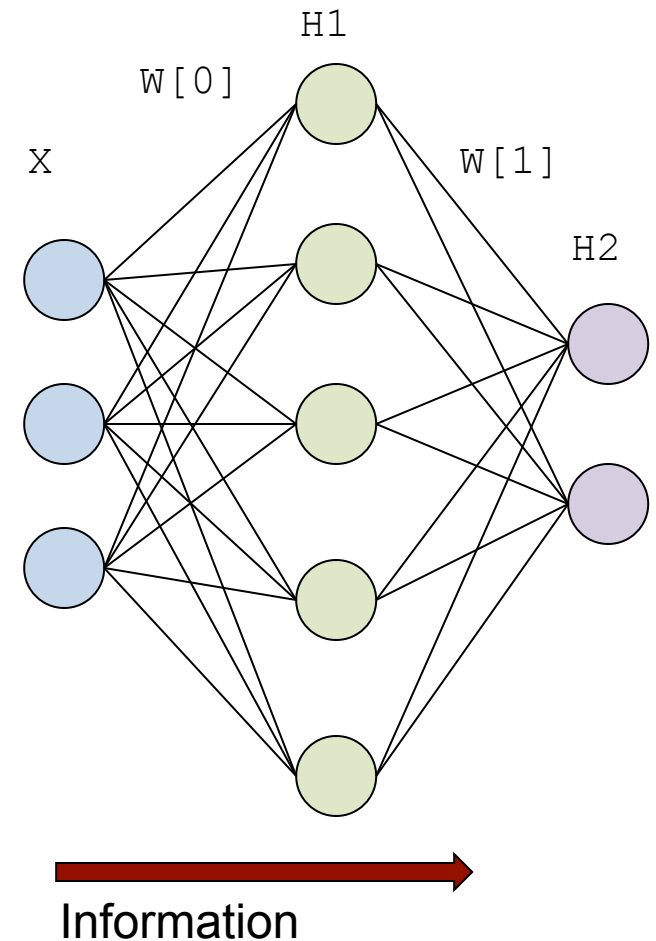
- Information flows left-to-right
  - Input observed features
  - Compute hidden nodes (parallel)
  - Compute next layer...

```
R = X.dot(W[0])+B[0]; # linear response
H1= Sig( R );         # activation f'n

S = H1.dot(W[1])+B[1]; # linear response
H2 = Sig( S );         # activation f'n

% ...
```

- Alternative: recurrent NNs...

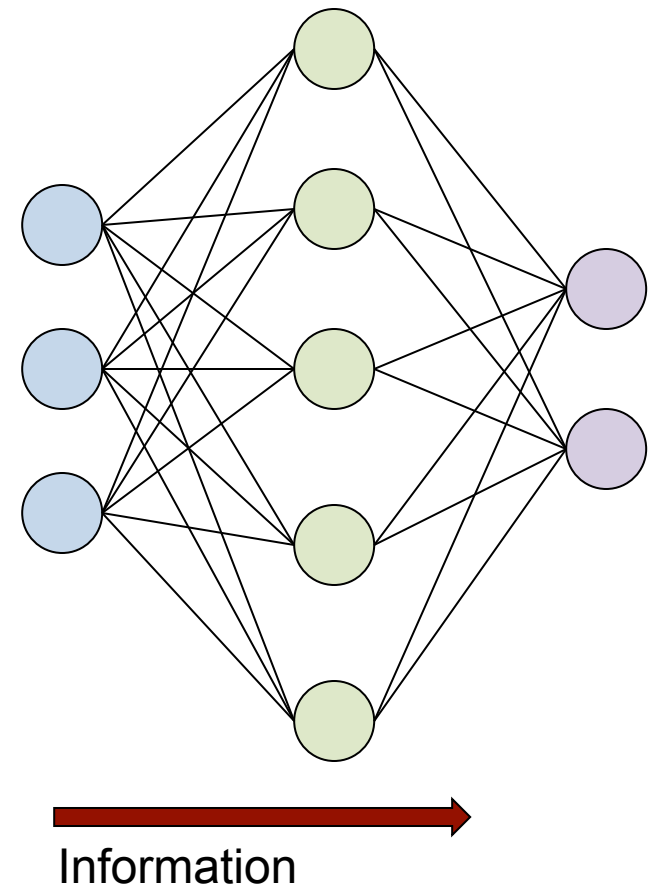




# Feed-forward networks

A note on multiple outputs:

- Regression:
  - Predict multi-dimensional  $y$
  - “Shared” representation  
= fewer parameters
- Classification
  - Predict binary vector
  - Multi-class classification  
 $y = 2 = [0 \ 0 \ 1 \ 0 \ \dots]$
  - Multiple, joint binary predictions  
(image tagging, etc.)
  - Often trained as regression (MSE),  
with saturating activation



+

# Machine Learning and Data Mining

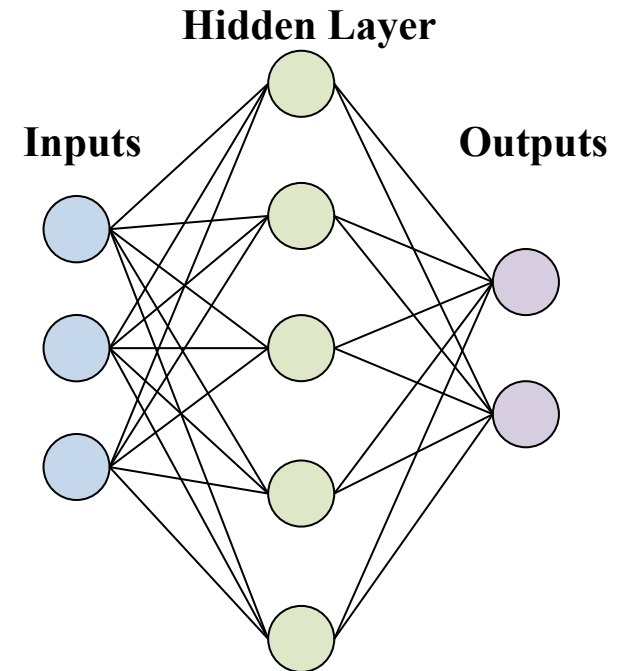
## Multi-layer Perceptrons & Neural Networks: Backpropagation

Prof. Alexander Ihler



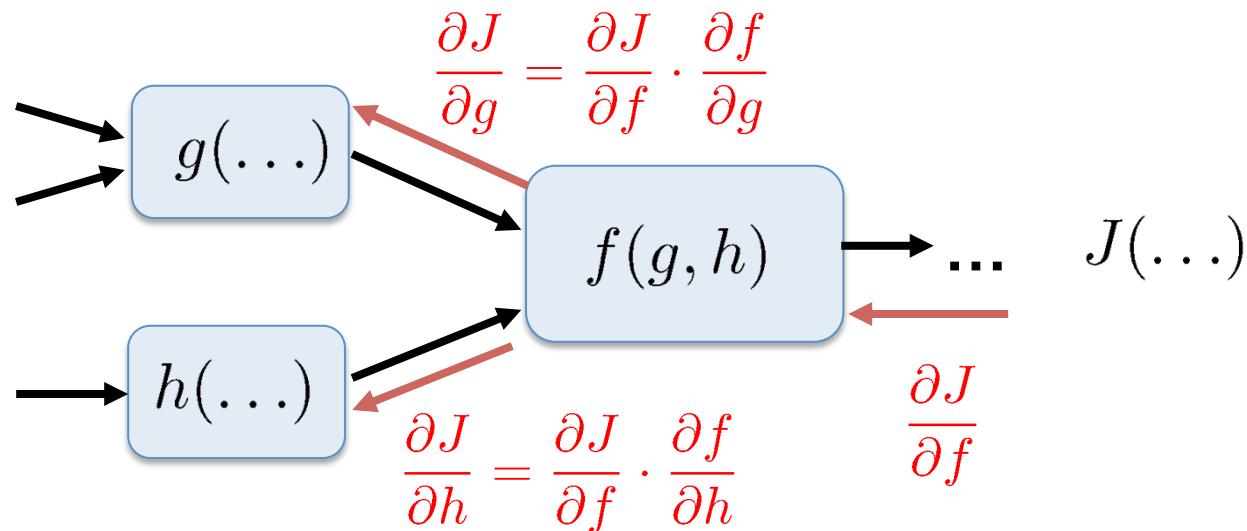
# Training MLPs

- Observe features “x” with target “y”
  - Push “x” through NN = output is “ $\hat{y}$ ”
  - Error:  $(y - \hat{y})^2$  (Can use different loss functions if desired...)
  - How should we update the weights to improve?
- 
- Single layer
    - Logistic sigmoid function
    - Smooth, differentiable
  - Optimize using:
    - Batch gradient descent
    - Stochastic gradient descent



# Gradient calculations

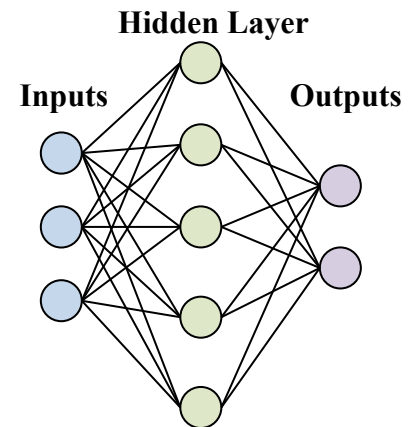
- Think of NNs as “schematics” made of smaller functions
  - Building blocks: summations & nonlinearities
  - For derivatives, just apply the chain rule, etc!



Ex:  $f(g, h) = g^2 h$

$$\frac{\partial J}{\partial g} = \frac{\partial J}{\partial f} \cdot 2g(\cdot)h(\cdot) \quad \frac{\partial J}{\partial h} = \frac{\partial J}{\partial f} \cdot g^2(\cdot)$$

save & reuse info (g,h) from forward computation!



# Backpropagation

- Just gradient descent...
- Apply the chain rule to the MLP

$$\begin{aligned}\frac{\partial J}{\partial w_{kj}^2} &= -2 \sum_{k'} (y_{k'} - \hat{y}_{k'}) (\partial \hat{y}_{k'}) \\ &= -2(y_k - \hat{y}_k) \sigma'(s_k) h_j\end{aligned}$$

(Identical to logistic mse regression with inputs “h<sub>j</sub>”)

## Forward pass

Loss function

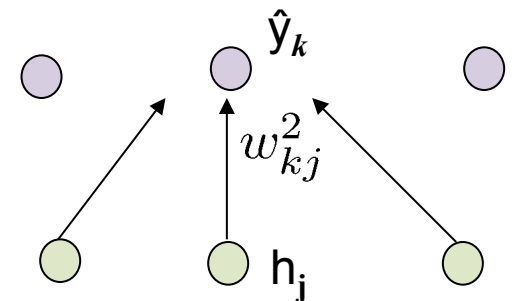
$$J_i(W) = \sum_k (y_k^{(i)} - \hat{y}_k^{(i)})^2$$

Output layer

$$\hat{y}_k = \sigma(s_k) = \sigma(\sum_j w_{kj}^2 h_j)$$

Hidden layer

$$h_j = \sigma(t_j) = \sigma(\sum_i w_{ji}^1 x_i)$$



# Backpropagation

- Just gradient descent...
- Apply the chain rule to the MLP

$$\frac{\partial J}{\partial w_{kj}^2} = -2 \sum_{k'} (y_{k'} - \hat{y}_{k'}) (\partial \hat{y}_{k'})$$

$$= \boxed{-2(y_k - \hat{y}_k) \sigma'(s_k)} h_j \quad (\text{Identical to logistic mse regression with inputs "h}_j\text{"})$$

$\beta_k^2$

$$\frac{\partial J}{\partial w_{ji}^1} = \sum_k -2(y_k - \hat{y}_k) (\partial \hat{y}_k)$$

$$= \sum_k -2(y_k - \hat{y}_k) \sigma'(s_k) w_{kj}^2 \partial h_j$$

$$= \sum_k \boxed{-2(y_k - \hat{y}_k) \sigma'(s_k)} w_{kj}^2 \sigma'(t_j) x_i$$

$\beta_k^2$

(c) Alexander Ihler

## Forward pass

Loss function

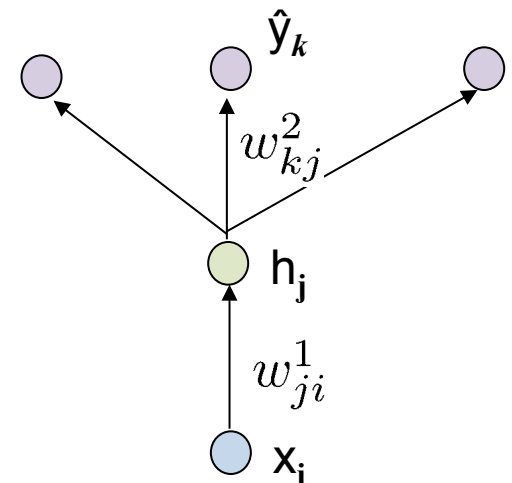
$$J_i(W) = \sum_k (y_k^{(i)} - \hat{y}_k^{(i)})^2$$

Output layer

$$\hat{y}_k = \sigma(s_k) = \sigma(\sum_j w_{kj}^2 h_j)$$

Hidden layer

$$h_j = \sigma(t_j) = \sigma(\sum_i w_{ji}^1 x_i)$$



# Backpropagation

- Just gradient descent...
- Apply the chain rule to the MLP

$$\frac{\partial J}{\partial w_{kj}^2} = -2(y_k - \hat{y}_k) \sigma'(s_k) h_j$$

$$\frac{\partial J}{\partial w_{ji}^1} = \sum_k \beta_k^2 -2(y_k - \hat{y}_k) \sigma'(s_k) w_{kj}^2 \sigma'(t_j) x_i$$

## Forward pass

Loss function

$$J_i(W) = \sum_k (y_k^{(i)} - \hat{y}_k^{(i)})^2$$

Output layer

$$\hat{y}_k = \sigma(s_k) = \sigma(\sum_j w_{kj}^2 h_j)$$

Hidden layer

$$h_j = \sigma(t_j) = \sigma(\sum_i w_{ji}^1 x_i)$$

```
% X : (1xN1)
H = Sig(X1.dot(W[0]))
% W1 : (N2 x N1+1)
% H : (1xN2)
Yh = Sig(H1.dot(W[1]))
% W2 : (N3 x N2+1)
% Yh : (1xN3)
```

```
B2 = (Y-Yhat) * dSig(S) # (1xN3)

G2 = B2.T.dot( H ) # (N3x1) * (1xN2) = (N3xN2)

B1 = B2.dot(W[1])*dSig(T) # (1xN3) . (N3*N2) * (1xN2)

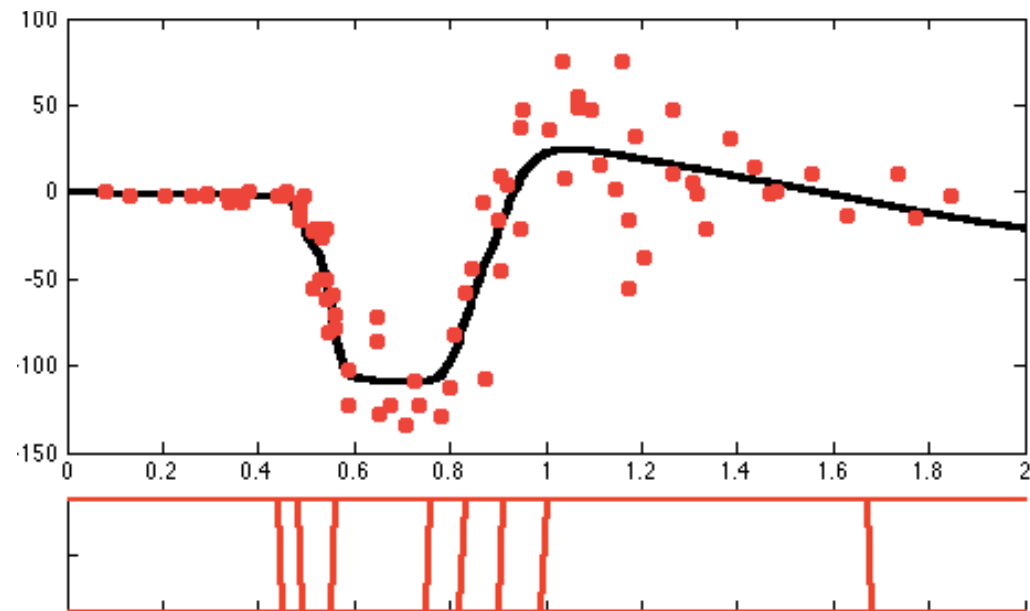
G1 = B1.T.dot( X ) # (N2 x N1+1)
```

# Example: Regression, MCycle data

- Train NN model, 2 layer
  - 1 input features => 1 input units
  - 10 hidden units
  - 1 target => 1 output units
  - Logistic sigmoid activation for hidden layer, linear for output layer

**Data:**  
+  
**learned prediction  $f'$ n:**

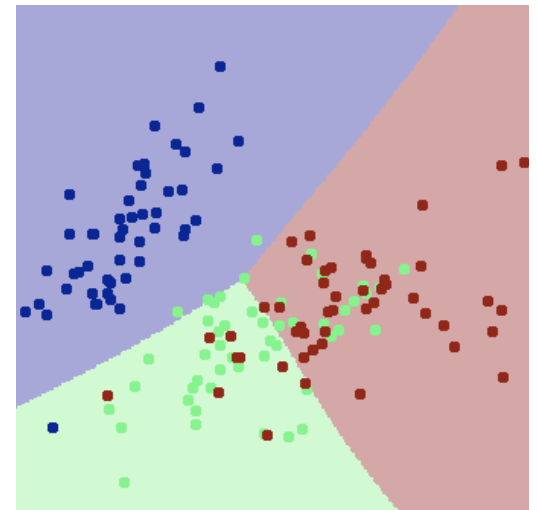
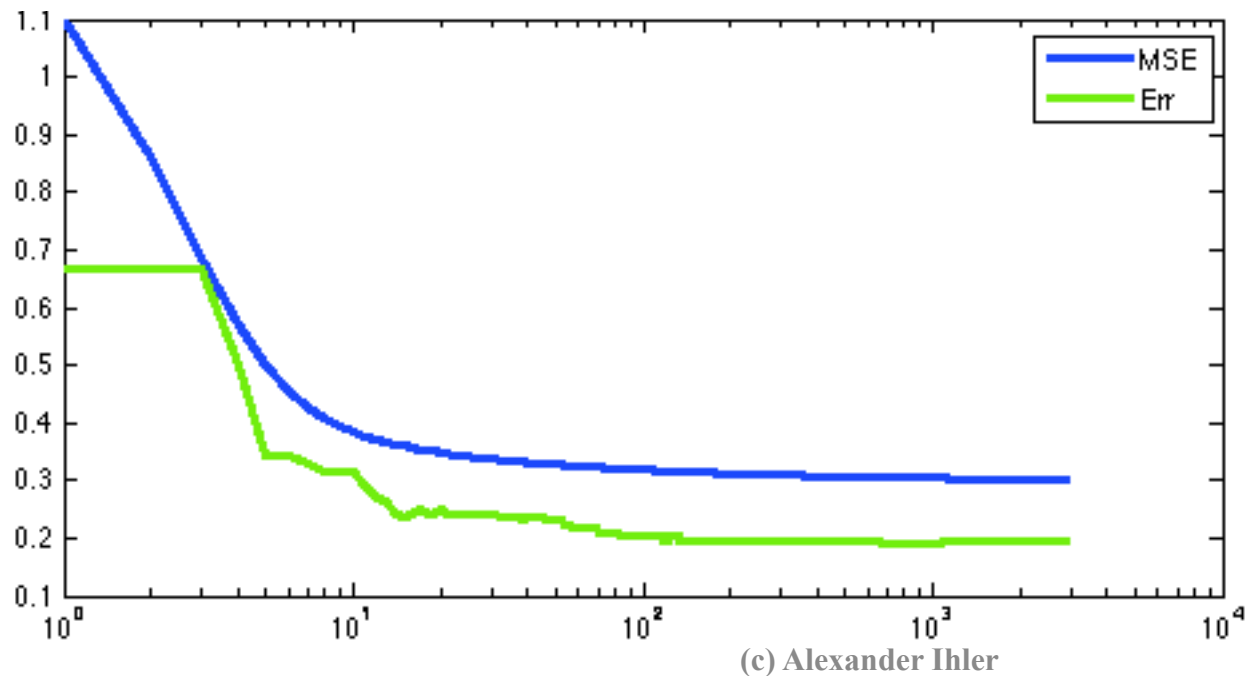
Responses of hidden nodes  
(= features of linear regression):  
select out useful regions of "x"





# Example: Classification, Iris data

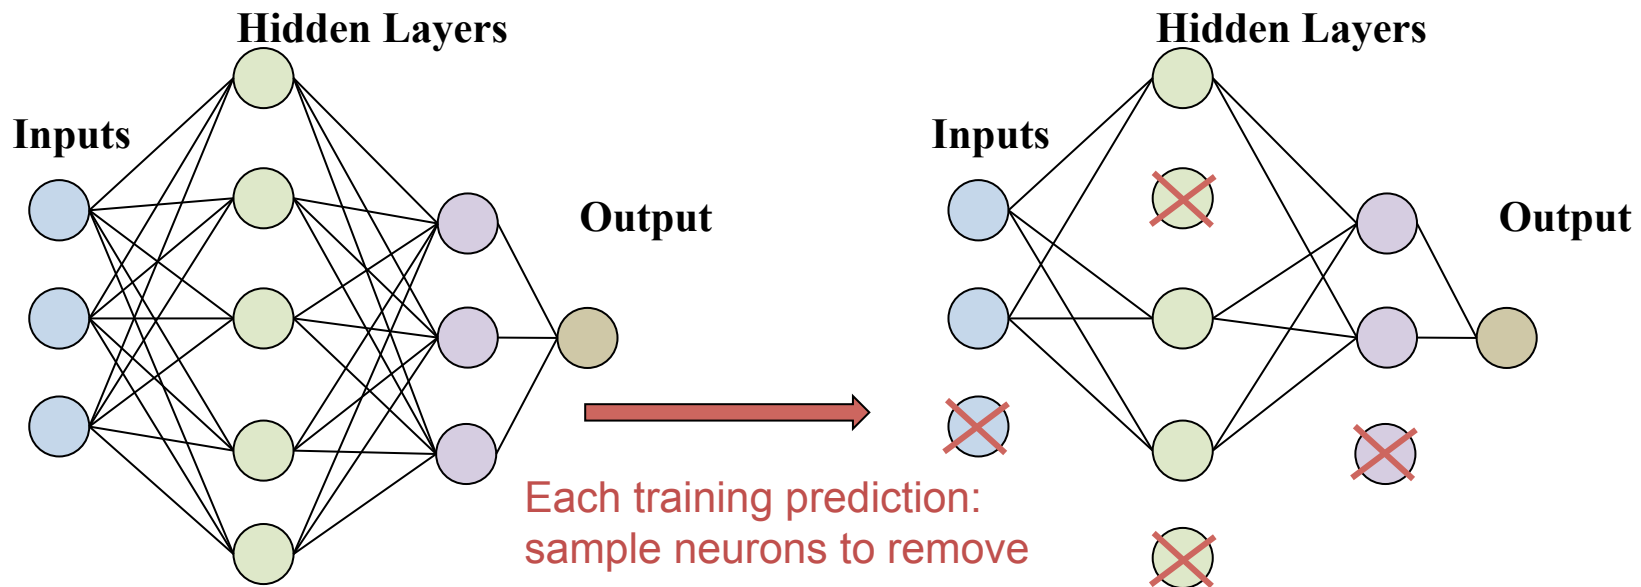
- Train NN model, 2 layer
  - 2 input features => 2 input units
  - 10 hidden units
  - 3 classes => 3 output units ( $y = [0 \ 0 \ 1]$ , etc.)
  - Logistic sigmoid activation functions
  - Optimize MSE of predictions using stochastic gradient



# Dropout

[Srivastava et al 2014]

- Another recent technique
  - Randomly “block” some neurons at each step
  - Trains model to have redundancy (predictions must be robust to blocking)



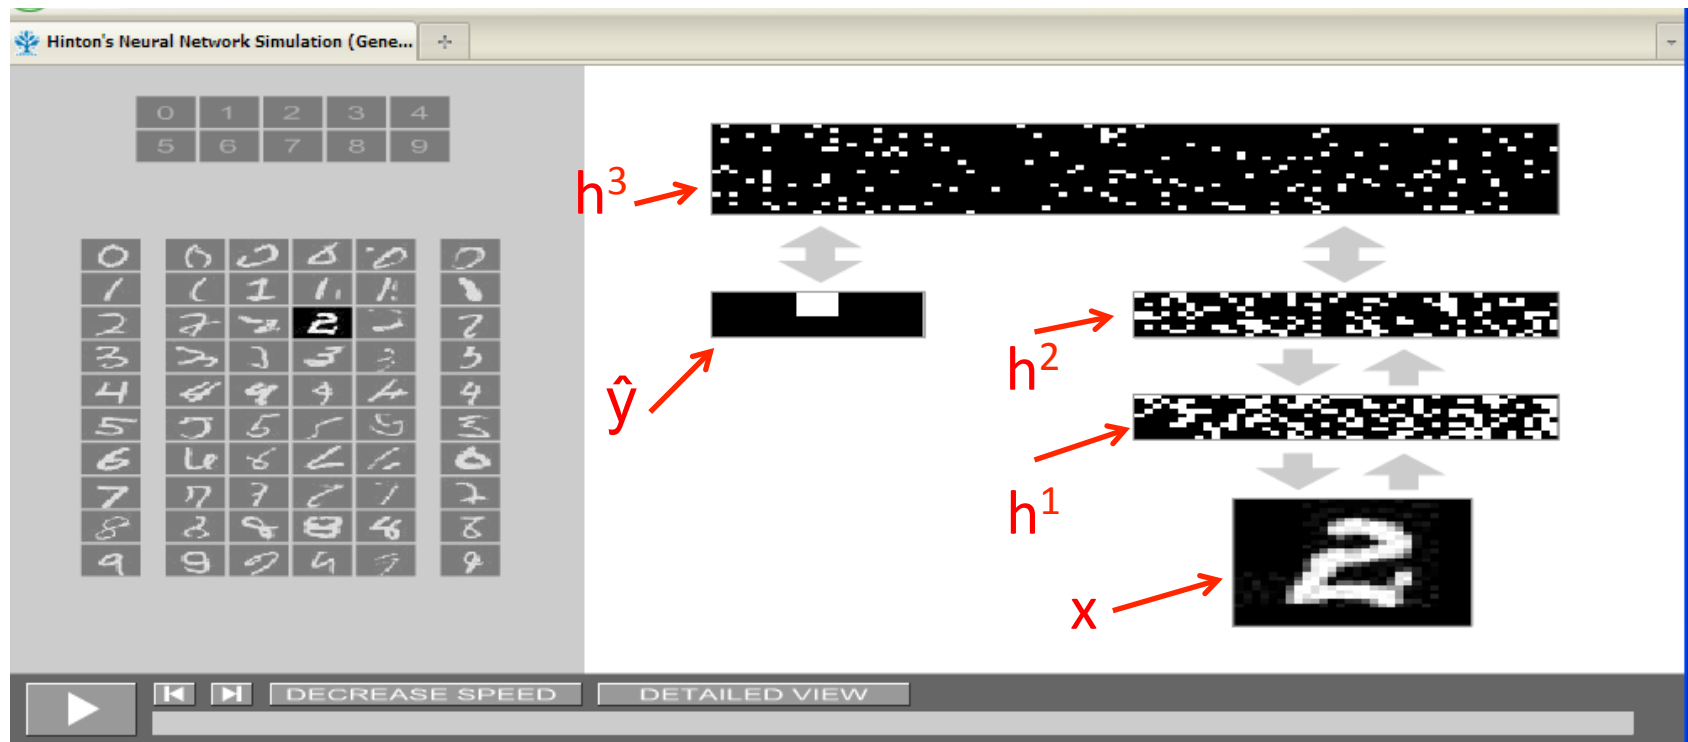
```
% ... during training ...  
R = X.dot(W[0])+B[0];           # linear response  
H1= Sig( R );                   # activation f'n  
H1 *= np.random.rand(*H1.shape)<p; #drop out!  
% ...
```

# MLPs in practice

[Hinton et al. 2007]

- Example: Deep belief nets
  - Handwriting recognition
  - Online demo
  - 784 pixels  $\Leftrightarrow$  500 mid  $\Leftrightarrow$  500 high  $\Leftrightarrow$  2000 top  $\Leftrightarrow$  10 labels

$x$        $h^1$        $h^2$        $h^3$        $\hat{y}$



(c) Alexander Ihler

# MLPs in practice

[Hinton et al. 2007]

- Example: Deep belief nets
  - Handwriting recognition
  - Online demo
  - 784 pixels  $\Leftrightarrow$  500 mid  $\Leftrightarrow$  500 high  $\Leftrightarrow$  2000 top  $\Leftrightarrow$  10 labels

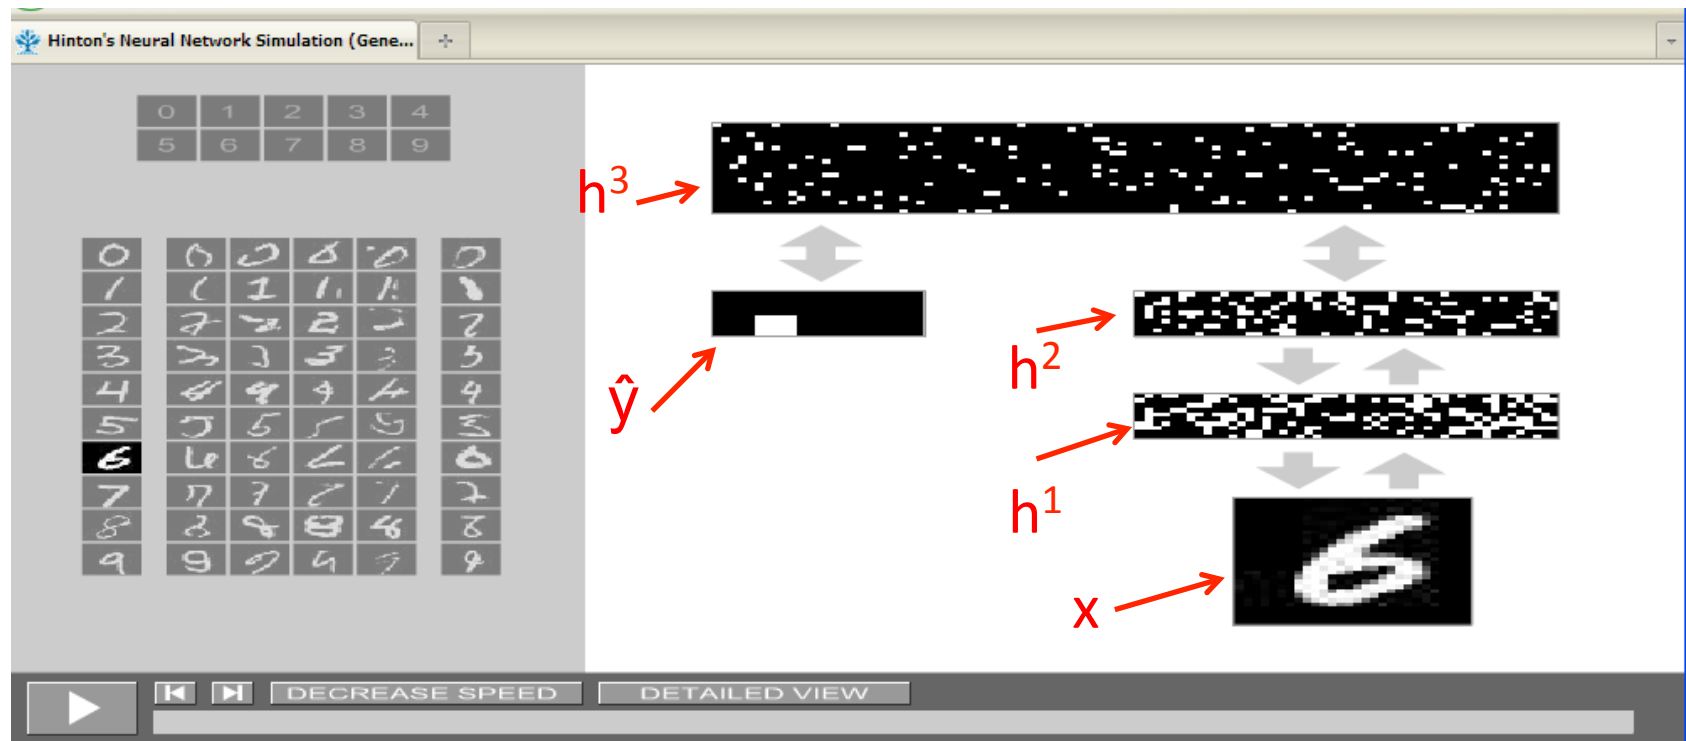
$x$

$h^1$

$h^2$

$h^3$

$\hat{y}$



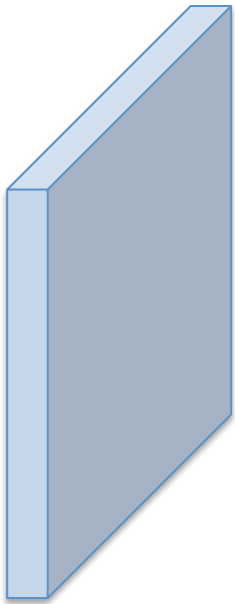
(c) Alexander Ihler

# Convolutional networks

- Organize & share the NN's weights (vs “dense”)
- Group weights into “filters”

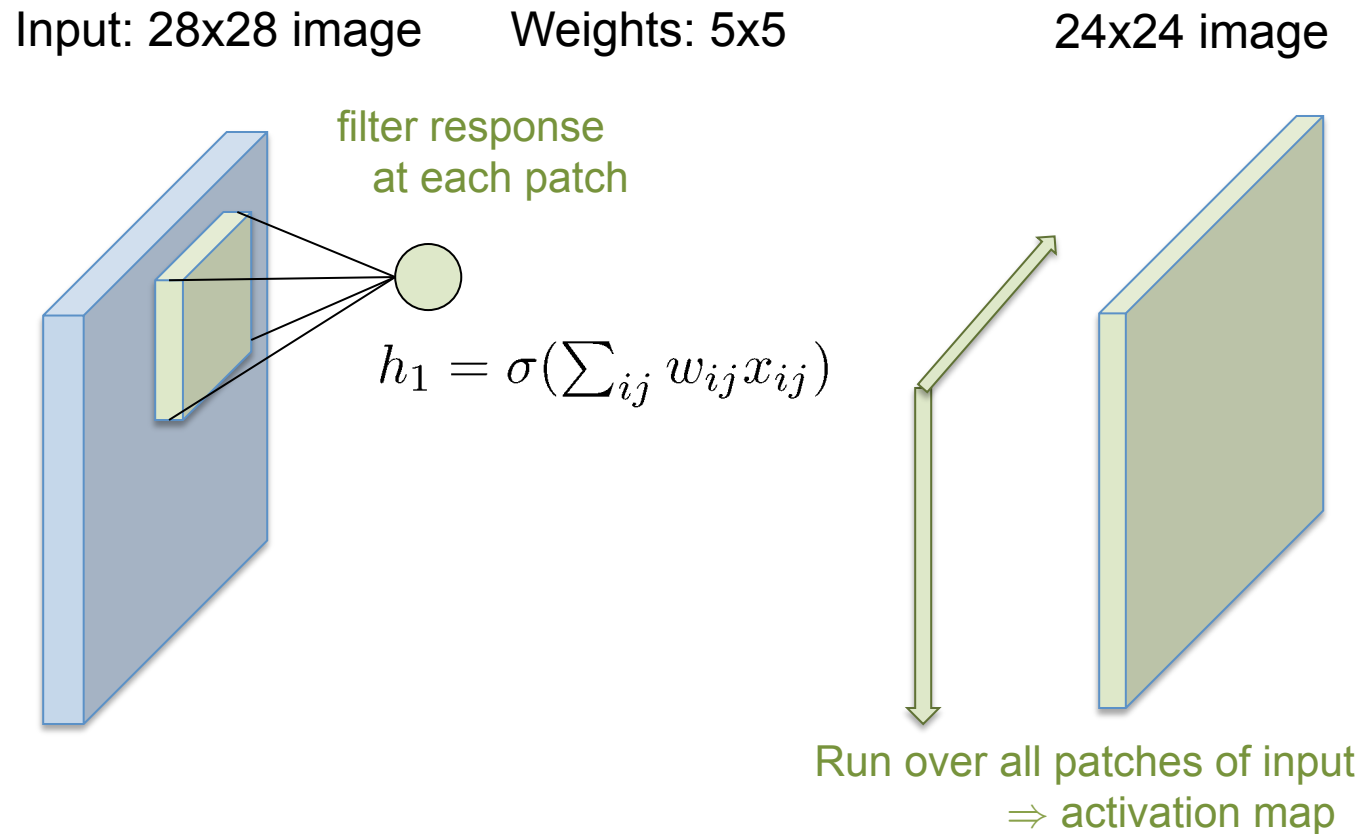
Input: 28x28 image

Weights: 5x5



# Convolutional networks

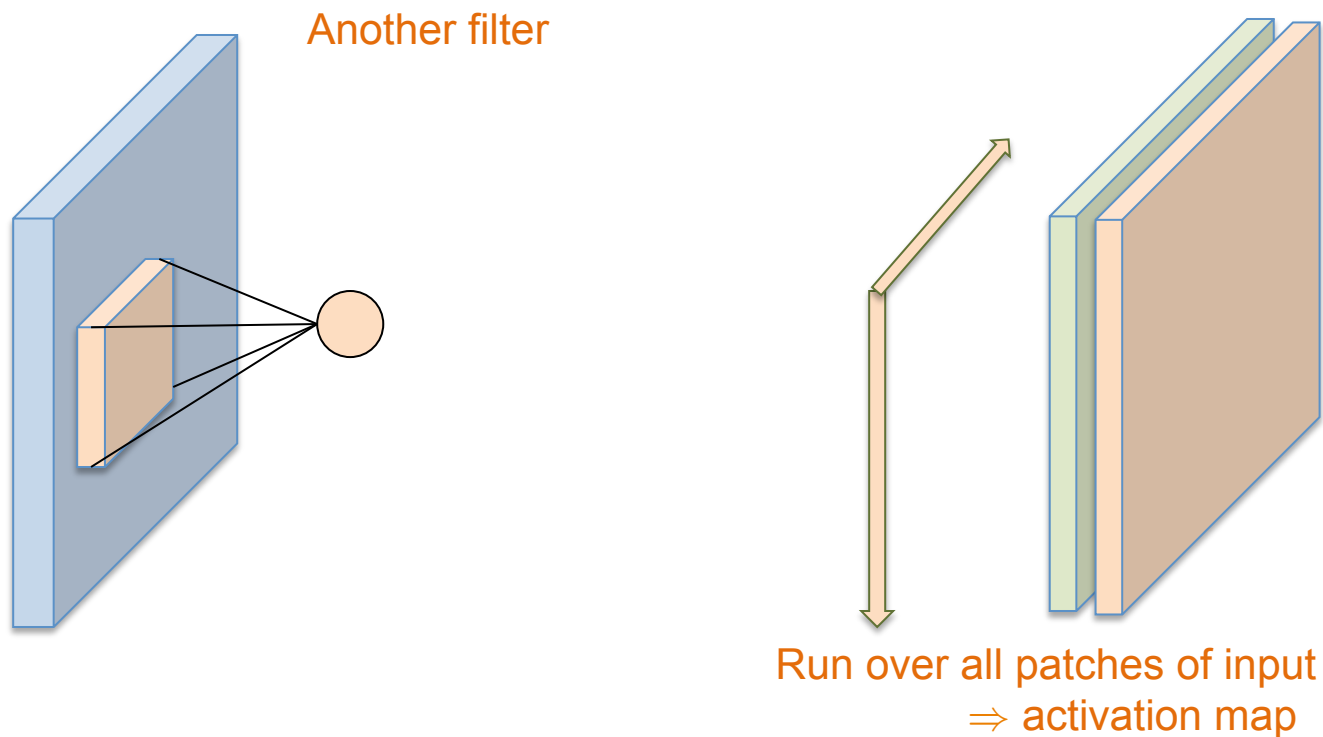
- Organize & share the NN's weights (vs “dense”)
- Group weights into “filters” & convolve across input image



# Convolutional networks

- Organize & share the NN's weights (vs “dense”)
- Group weights into “filters” & convolve across input image

Input: 28x28 image      Weights: 5x5



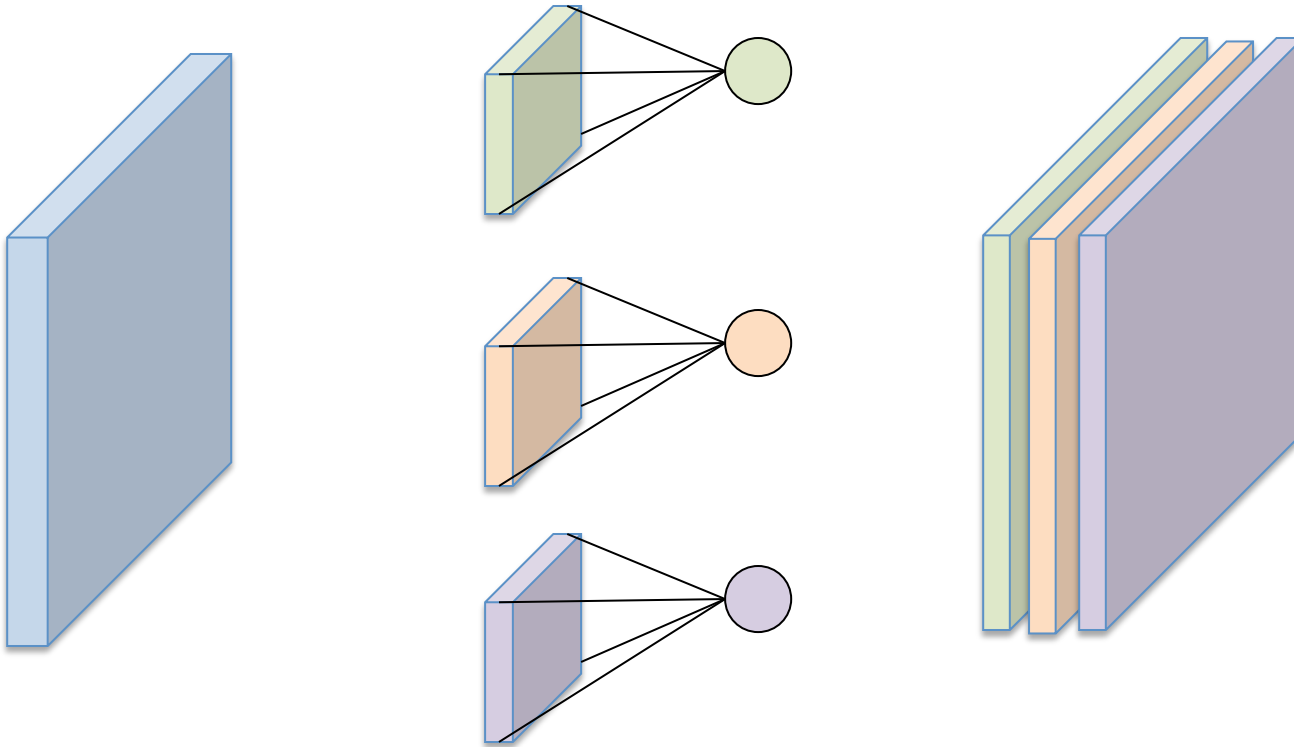
# Convolutional networks

- Organize & share the NN's weights (vs “dense”)
- Group weights into “filters” & convolve across input image
- Many hidden nodes, but few parameters!

Input: 28x28 image

Weights: 5x5

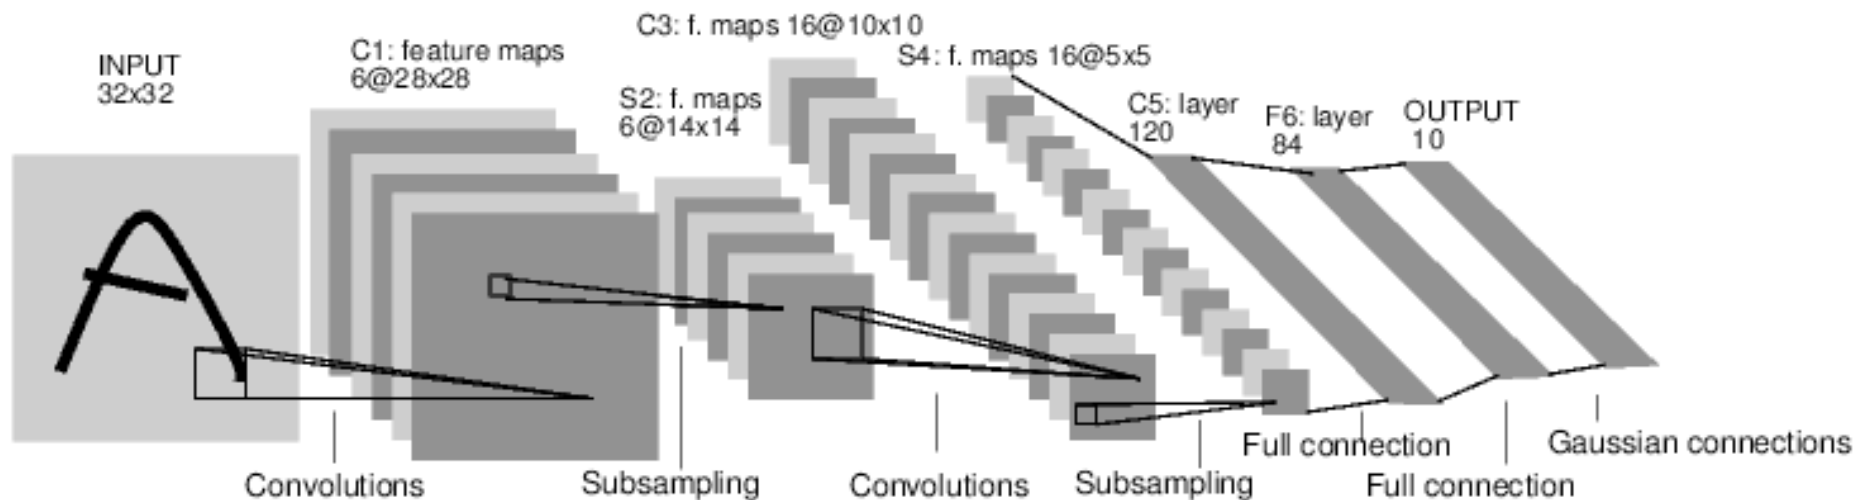
Hidden layer 1





# Convolutional networks

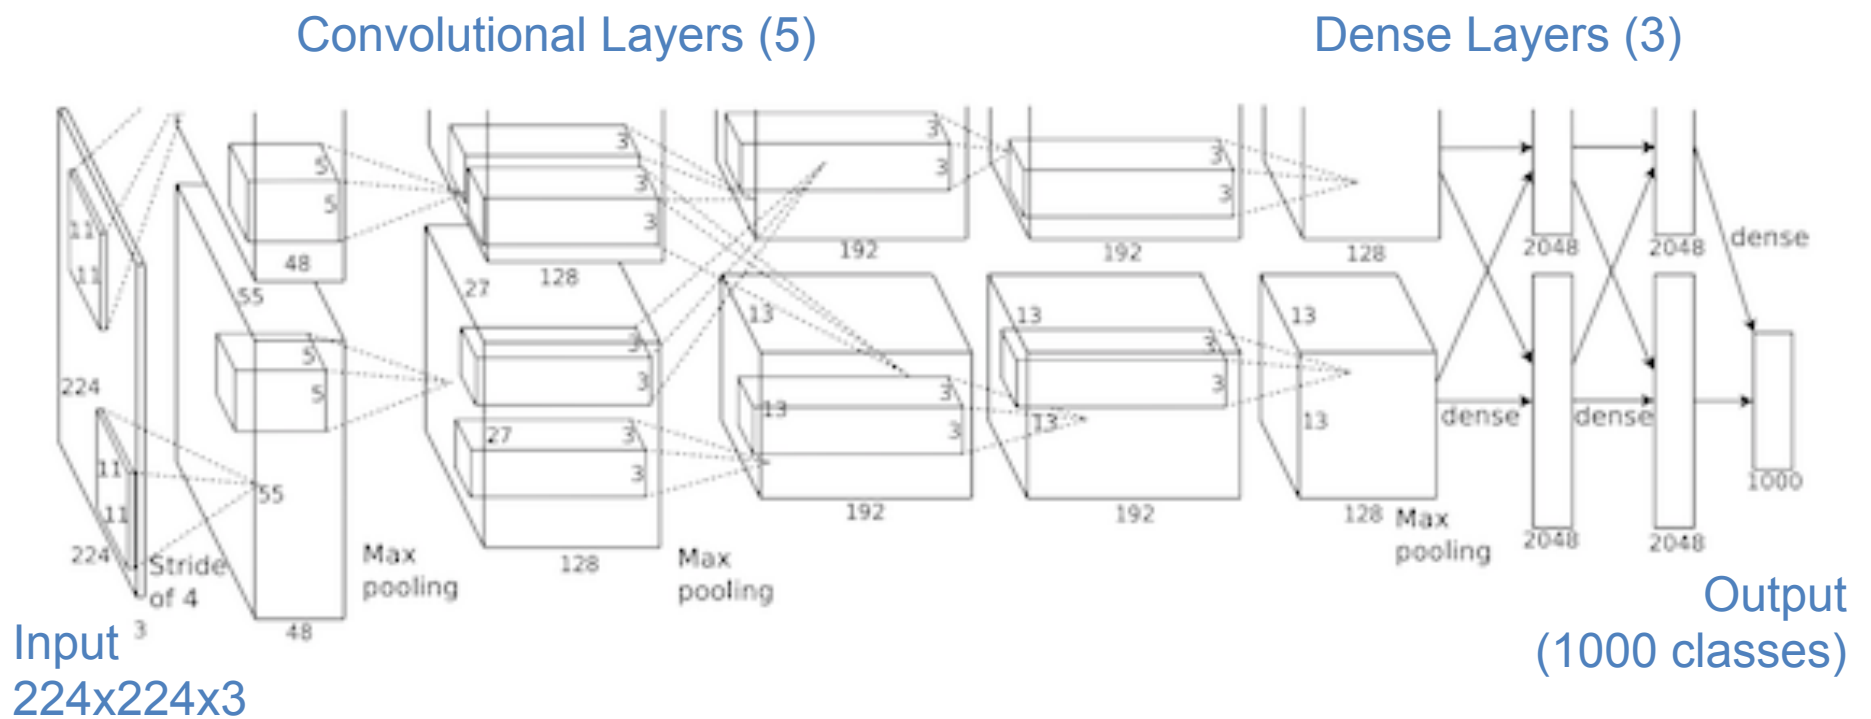
- Again, can view components as building blocks
- Design overall, deep structure from parts
  - Convolutional layers
  - “Max-pooling” (sub-sampling) layers
  - Densely connected layers



LeNet-5 [LeCun 1980]

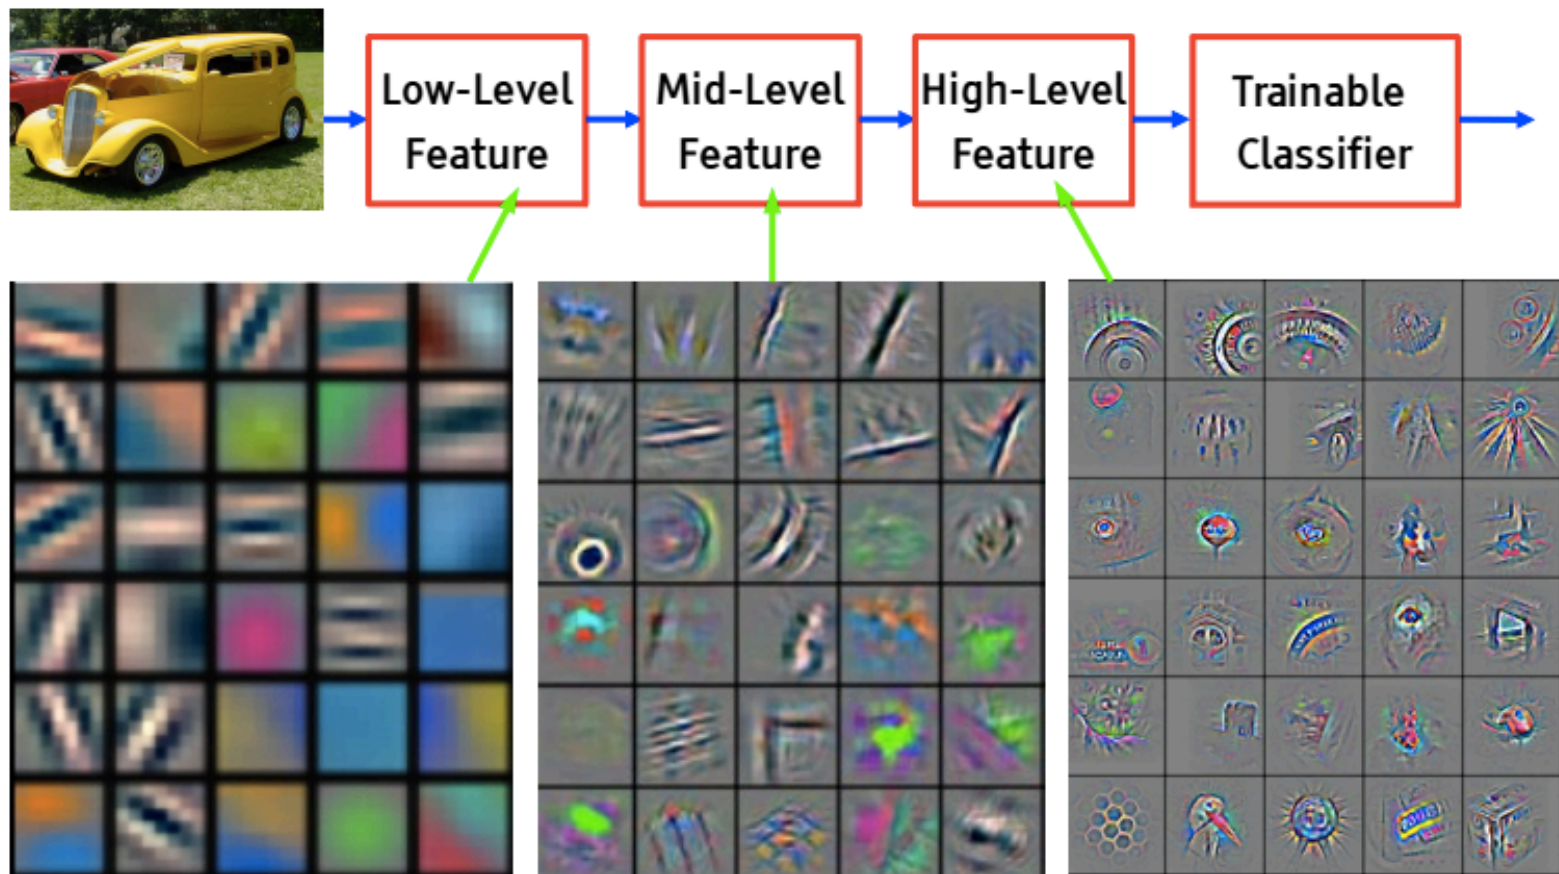
# Ex: AlexNet

- Deep NN model for ImageNet classification
  - 650k units; 60m parameters
  - 1m data; 1 week training (GPUs)



# Hidden layers as “features”

- Visualizing a convolutional network’s filters [Zeiler & Fergus 2013]



Slide image from Yann LeCun:  
<https://drive.google.com/open?id=0BxKBnD5y2M8NcIFWSXNxa0JIZTg>

# Neural networks & DBNs

- Want to try them out?
- Matlab “Deep Learning Toolbox”  
<https://github.com/rasmusbergpalm/DeepLearnToolbox>



rasmusbergpalm / DeepLearnToolbox

---

Matlab/Octave toolbox for deep learning. Includes Deep Belief Nets, Stacked Autoencoders, Convolutional Neural Nets, Convolutional Autoencoders and vanilla Neural Nets. Each method has examples to get you started.

- PyLearn2  
<https://github.com/lisa-lab/pylearn2>
- TensorFlow

# Summary

- Neural networks, multi-layer perceptrons
- Cascade of simple perceptrons
  - Each just a linear classifier
  - Hidden units used to create new features
- Together, general function approximators
  - Enough hidden units (features) = any function
  - Can create nonlinear classifiers
  - Also used for function approximation, regression, ...
- Training via backprop
  - Gradient descent; logistic; apply chain rule. Building block view.
- Advanced: deep nets, conv nets, dropout, ...