

+

Machine Learning and Data Mining

Decision Trees

Prof. Alexander Ihler




Decision trees

- Functional form $f(x;\theta)$: nested “if-then-else” statements
 - Discrete features: fully expressive (any function)
- Structure:
 - Internal nodes: check feature, branch on value
 - Leaf nodes: output prediction

“XOR”

x_1	x_2	y
0	0	1
0	1	-1
1	0	-1
1	1	1





```
if X1:                                     # branch on feature at root
    if X2: return +1                       # if true, branch on right child feature
    else:  return -1                       # & return leaf value
else: # left branch:
    if X2: return -1                       # branch on left child feature
    else:  return +1                       # & return leaf value
```

Parameters?

Tree structure, features, and leaf outputs


Decision trees

- Real-valued features
 - Compare feature value to some threshold





Decision trees

- Categorical variables
 - Could have one child per value
 - Binary splits: single values, or by subsets



The discrete variable will not appear again below here...





Could appear again multiple times...

(This ^^ is easy to implement using a 1-of-K representation...)


Decision trees

- “Complexity” of function depends on the depth
- A depth-1 decision tree is called a decision “stump”
 - Simpler than a linear classifier!




Decision trees

- “Complexity” of function depends on the depth
- More splits provide a finer-grained partitioning




Depth d = up to 2^d regions & predictions




Decision trees for regression

- Exactly the same
- Predict real valued numbers at leaf nodes
- Examples on a single scalar feature:

Depth 1 = 2 regions & predictions



Depth 2 = 4 regions & predictions



...

Machine Learning and Data Mining

Learning Decision Trees

Prof. Alexander Ihler



Learning decision trees

- Break into two parts
 - Should this be a leaf node?
 - If so: what should we predict?
 - If not: how should we further split the data?
- Leaf nodes: best prediction given this data subset
 - Classify: pick majority class; Regress: predict average value
- Non-leaf nodes: pick a feature and a split
 - Greedy: “score” all possible features and splits
 - Score function measures “purity” of data after split
 - How much easier is our prediction task after we divide the data?
- When to make a leaf node?
 - All training examples the same class (correct), or indistinguishable
 - Fixed depth (fixed complexity decision boundary)
 - Others ...

Example algorithms:
ID3, C4.5

See e.g. wikipedia,
“Classification and
regression tree”

Learning decision trees

Algorithm 1 BuildTree(D): Greedy training of a decision tree

Input: A data set $D = (X, Y)$.

Output: A decision tree.

if LeafCondition(D) **then**

$f_n = \text{FindBestPrediction}(D)$

else

$j_n, t_n = \text{FindBestSplit}(D)$

$D_L = \{(x^{(i)}, y^{(i)}) : x_{j_n}^{(i)} < t_n\}$ and

$D_R = \{(x^{(i)}, y^{(i)}) : x_{j_n}^{(i)} \geq t_n\}$



 leftChild = BuildTree(D_L)

 rightChild = BuildTree(D_R)

end if

Scoring decision tree splits

- Suppose we are considering splitting feature 1
 - How can we score any particular split?
 - “Impurity” – how easy is the prediction problem in the leaves?
- “Greedy” – could choose split with the best accuracy
 - Assume we have to predict a value next
 - MSE (regression)
 - 0/1 loss (classification)
- But: “soft” score can work better




Entropy and information

- “Entropy” is a measure of randomness
 - How hard is it to communicate a result to you?
 - Depends on the probability of the outcomes
- Communicating fair coin tosses
 - Output: H H T H T T T H H H H T ...
 - Sequence takes n bits – each outcome totally unpredictable
- Communicating my daily lottery results
 - Output: 0 0 0 0 0 0 ...
 - Most likely to take one bit – I lost every day.
 - Small chance I’ll have to send more bits (won & when)

Lost: 0
Won 1: 1(...)0
Won 2: 1(...)1(...)0
- Takes less work to communicate because it’s less random
 - Use a few bits for the most likely outcome, more for less likely ones


Entropy and information

- Entropy $H(x) \equiv \mathbb{E}[\log 1/p(x)] = \sum p(x) \log 1/p(x)$
 - Log base two, units of entropy are “bits”
 - Two outcomes: $H = -p \log(p) - (1-p) \log(1-p)$
- Examples:




$$\begin{aligned} H(x) &= .25 \log 4 + .25 \log 4 + \\ &\quad .25 \log 4 + .25 \log 4 \\ &= \log 4 = 2 \text{ bits} \end{aligned}$$

Max entropy for 4 outcomes



$$\begin{aligned} H(x) &= .75 \log 4/3 + .25 \log 4 \\ &\approx .8133 \text{ bits} \end{aligned}$$







$$\begin{aligned} H(x) &= 1 \log 1 \\ &= 0 \text{ bits} \end{aligned}$$

Min entropy

Entropy and information

- Information gain
 - How much is entropy reduced by measurement?
- Information: expected information gain







$$\text{Information} = 13/18 * (.99-.77) + 5/18 * (.99 - 0)$$

$$\text{Equivalent: } \sum p(s,c) \log [p(s,c) / p(s) p(c)]$$

$$= 10/18 \log[(10/18) / (13/18) (10/18)] + 3/18 \log[(3/18)/(13/18)(8/18)] + \dots$$

Entropy and information

- Information gain
 - How much is entropy reduced by measurement?
- Information: expected information gain







$$\text{Information} = 17/18 * (.99-.97) + 1/18 * (.99 - 0)$$

Less information reduction – a less desirable split of the data

Gini index & impurity


- An alternative to information gain
 - Measures variance in the allocation (instead of entropy)
- $H_{\text{gini}} = \sum_c p(c) (1-p(c))$ vs. $H_{\text{ent}} = - \sum_c p(c) \log p(c)$



$$\text{Gini Index} = 13/18 * (.494 - .355) + 5/18 * (.494 - 0)$$


Entropy vs Gini impurity

- The two are nearly the same...
 - Pick whichever one you like




For regression

- Most common is to measure variance reduction
 - Equivalent to “information gain” in a Gaussian model...



Var = .25



Var = .1
Prob = 4/10

Var = .2
Prob = 6/10

Var reduction = $4/10 * (.25 - .1) + 6/10 * (.25 - .2)$

Scoring decision tree splits

Algorithm 1 FindBestSplit(D)

Input: A data set $D = (X, Y)$ of size m ;
impurity function $H(\cdot)$.

Output: A split j^*, t^* minimizing impurity H

Initialize $H^* = 0$

for each feature j **do**

Sort $\{x_j^{(i)}\}$ in order of increasing value

for each i such that $x^{(i)} < x^{(i+1)}$ **do**

Compute $p_c^L = \frac{1}{i} \sum_{k \leq i} \mathbb{1}[y^{(k)} = c]$

and $p_c^R = \frac{1}{k-i} \sum_{k > i} \mathbb{1}[y^{(k)} = c]$

Set $H' = \frac{i}{m} H(p^L) + \frac{m-i}{m} H(p^R)$

if $H' < H^*$ **then**

Set $j^* = j, t^* = (x^{(i)} - x^{(i+1)})/2, H^* = H'$

end if

end for

end for

Return j^*, t^*


Building a decision tree

Algorithm 1 BuildTree(D): Greedy training of a decision tree

Input: A data set $D = (X, Y)$.

Output: A decision tree.

```
if LeafCondition( $D$ ) then
     $f_n = \text{FindBestPrediction}(D)$ 
else
     $j_n, t_n = \text{FindBestSplit}(D)$ 
     $D_L = \{(x^{(i)}, y^{(i)}) : x_{j_n}^{(i)} < t_n\}$     and
     $D_R = \{(x^{(i)}, y^{(i)}) : x_{j_n}^{(i)} \geq t_n\}$ 
    leftChild = BuildTree( $D_L$ )
    rightChild = BuildTree( $D_R$ )
end if
```




Stopping conditions:

- * # of data $< K$
- * Depth $> D$
- * All data indistinguishable (discrete features)
- * Prediction sufficiently accurate


- * Information gain threshold?
Often not a good idea!
No single split improves,
but, two splits do.
Better: build full tree, then prune

Controlling complexity


- Maximum depth cutoff




No limit




Depth 1




Depth 2



Depth 3







Depth 4



Depth 5

Controlling complexity

- Minimum # parent data



- Alternate (similar): min # of data per leaf


Decision trees in python

- Many implementations
- Class implementation:
 - real-valued features (can use 1-of-k for discrete)
 - Uses entropy (easy to extend)

```
T = dt.treeClassify()  
T.train(X,Y,maxDepth=2)  
print T
```

```
if x[0] < 5.602476:  
    if x[1] < 3.009747:  
        Predict 1.0          # green  
    else:  
        Predict 0.0          # blue  
else:  
    if x[0] < 6.186588:  
        Predict 1.0          # green  
    else:  
        Predict 2.0          # red
```

```
ml.plotClassify2D(T, X,Y)
```



Summary

- Decision trees
 - Flexible functional form
 - At each level, pick a variable and split condition
 - At leaves, predict a value
- Learning decision trees
 - Score all splits & pick best
 - Classification: Information gain, Gini index
 - Regression: Expected variance reduction
 - Stopping criteria
- Complexity depends on depth
 - Decision stumps: very simple classifiers